



PostgreSQL

DBA Training

Surendra Panpaliya

Agenda



DAY 1 — PostgreSQL Foundations, Architecture & Tools



DAY 2 — Transactions, Locking & Maintenance



DAY 3 — Query Optimization & Indexing



DAY 4 — Backup, Replication Concepts & Migration

A. PostgreSQL Introduction

PostgreSQL capabilities overview

PostgreSQL vs Oracle vs SQL Server vs AWS RDS PostgreSQL

Architecture differences (Local vs RDS)

MVCC high-level concept

Key differences vs enterprise databases

A. PostgreSQL Introduction

Fidelity Use Case

PostgreSQL used for

Investment reporting, portfolio views, microservices

Oracle retained for **core settlement / legacy systems**

RDS PostgreSQL used for **managed production workloads**

B. PostgreSQL Local Architecture (1 hour)



Cluster structure



(data directory, instance, databases, schemas)



Key processes



postmaster, autovacuum, WAL writer, checkpointer

B. PostgreSQL Local Architecture (1 hour)

Memory architecture

WAL, checkpoints, background writer

MVCC row versioning

Fidelity Use Case

Gurugram supports multi

DB clusters → shared memory understanding critical

UK latency SLAs → WAL & checkpoint behavior matters

China read-heavy compliance queries → MVCC benefits

C. Tools Overview & Connectivity (30 minutes)



psql vs pgAdmin (DBA usage patterns)



Connecting locally

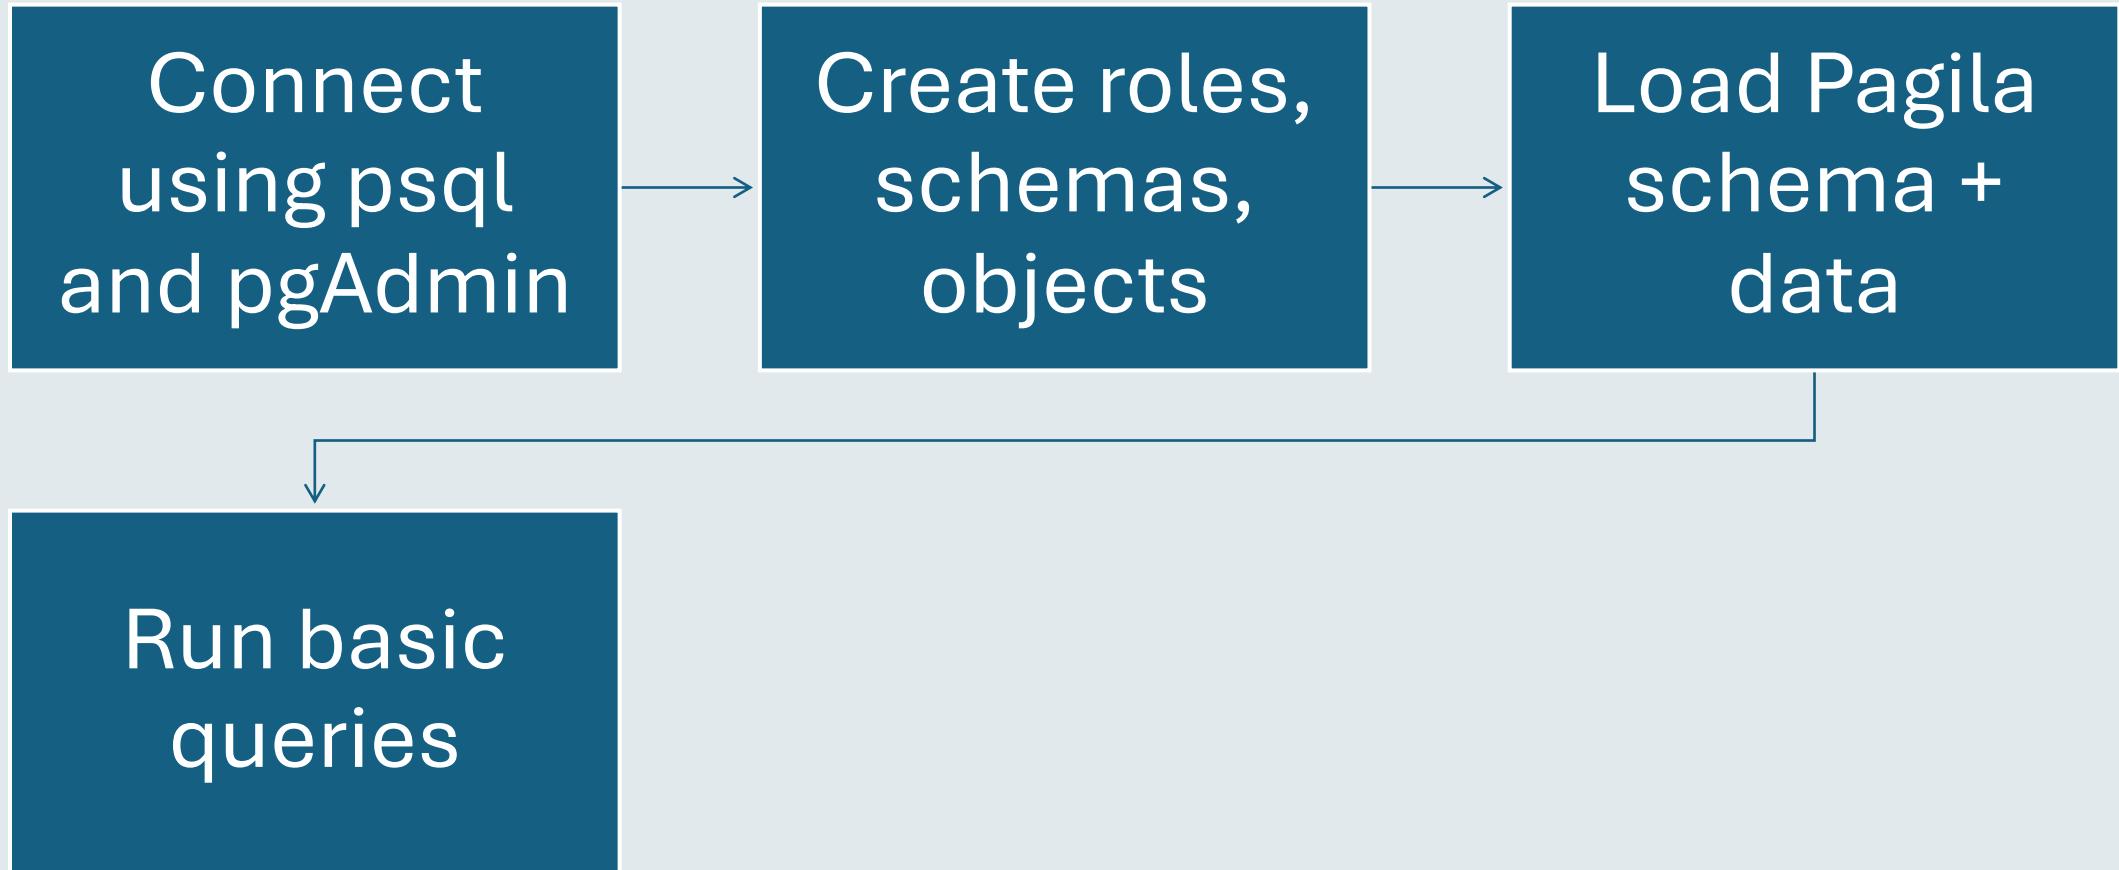


Roles, databases, schemas



Pagila dataset overview

D. Hands-On Lab



D. Hands-On Lab

Measure
performance
using:

\timing

EXPLAIN /
EXPLAIN
ANALYZE

Explore system
catalogs:

pg_class

pg_stat_activity

Module-wise FAQ — Day 1



Q1. Is PostgreSQL production-ready for financial systems?



Q2. How is PostgreSQL different from Oracle RAC?



Q3. Why do DBAs still need architecture knowledge in RDS?

Agenda



PostgreSQL vs Oracle: Key architectural differences



PostgreSQL installation & tools (pgAdmin, DBeaver)

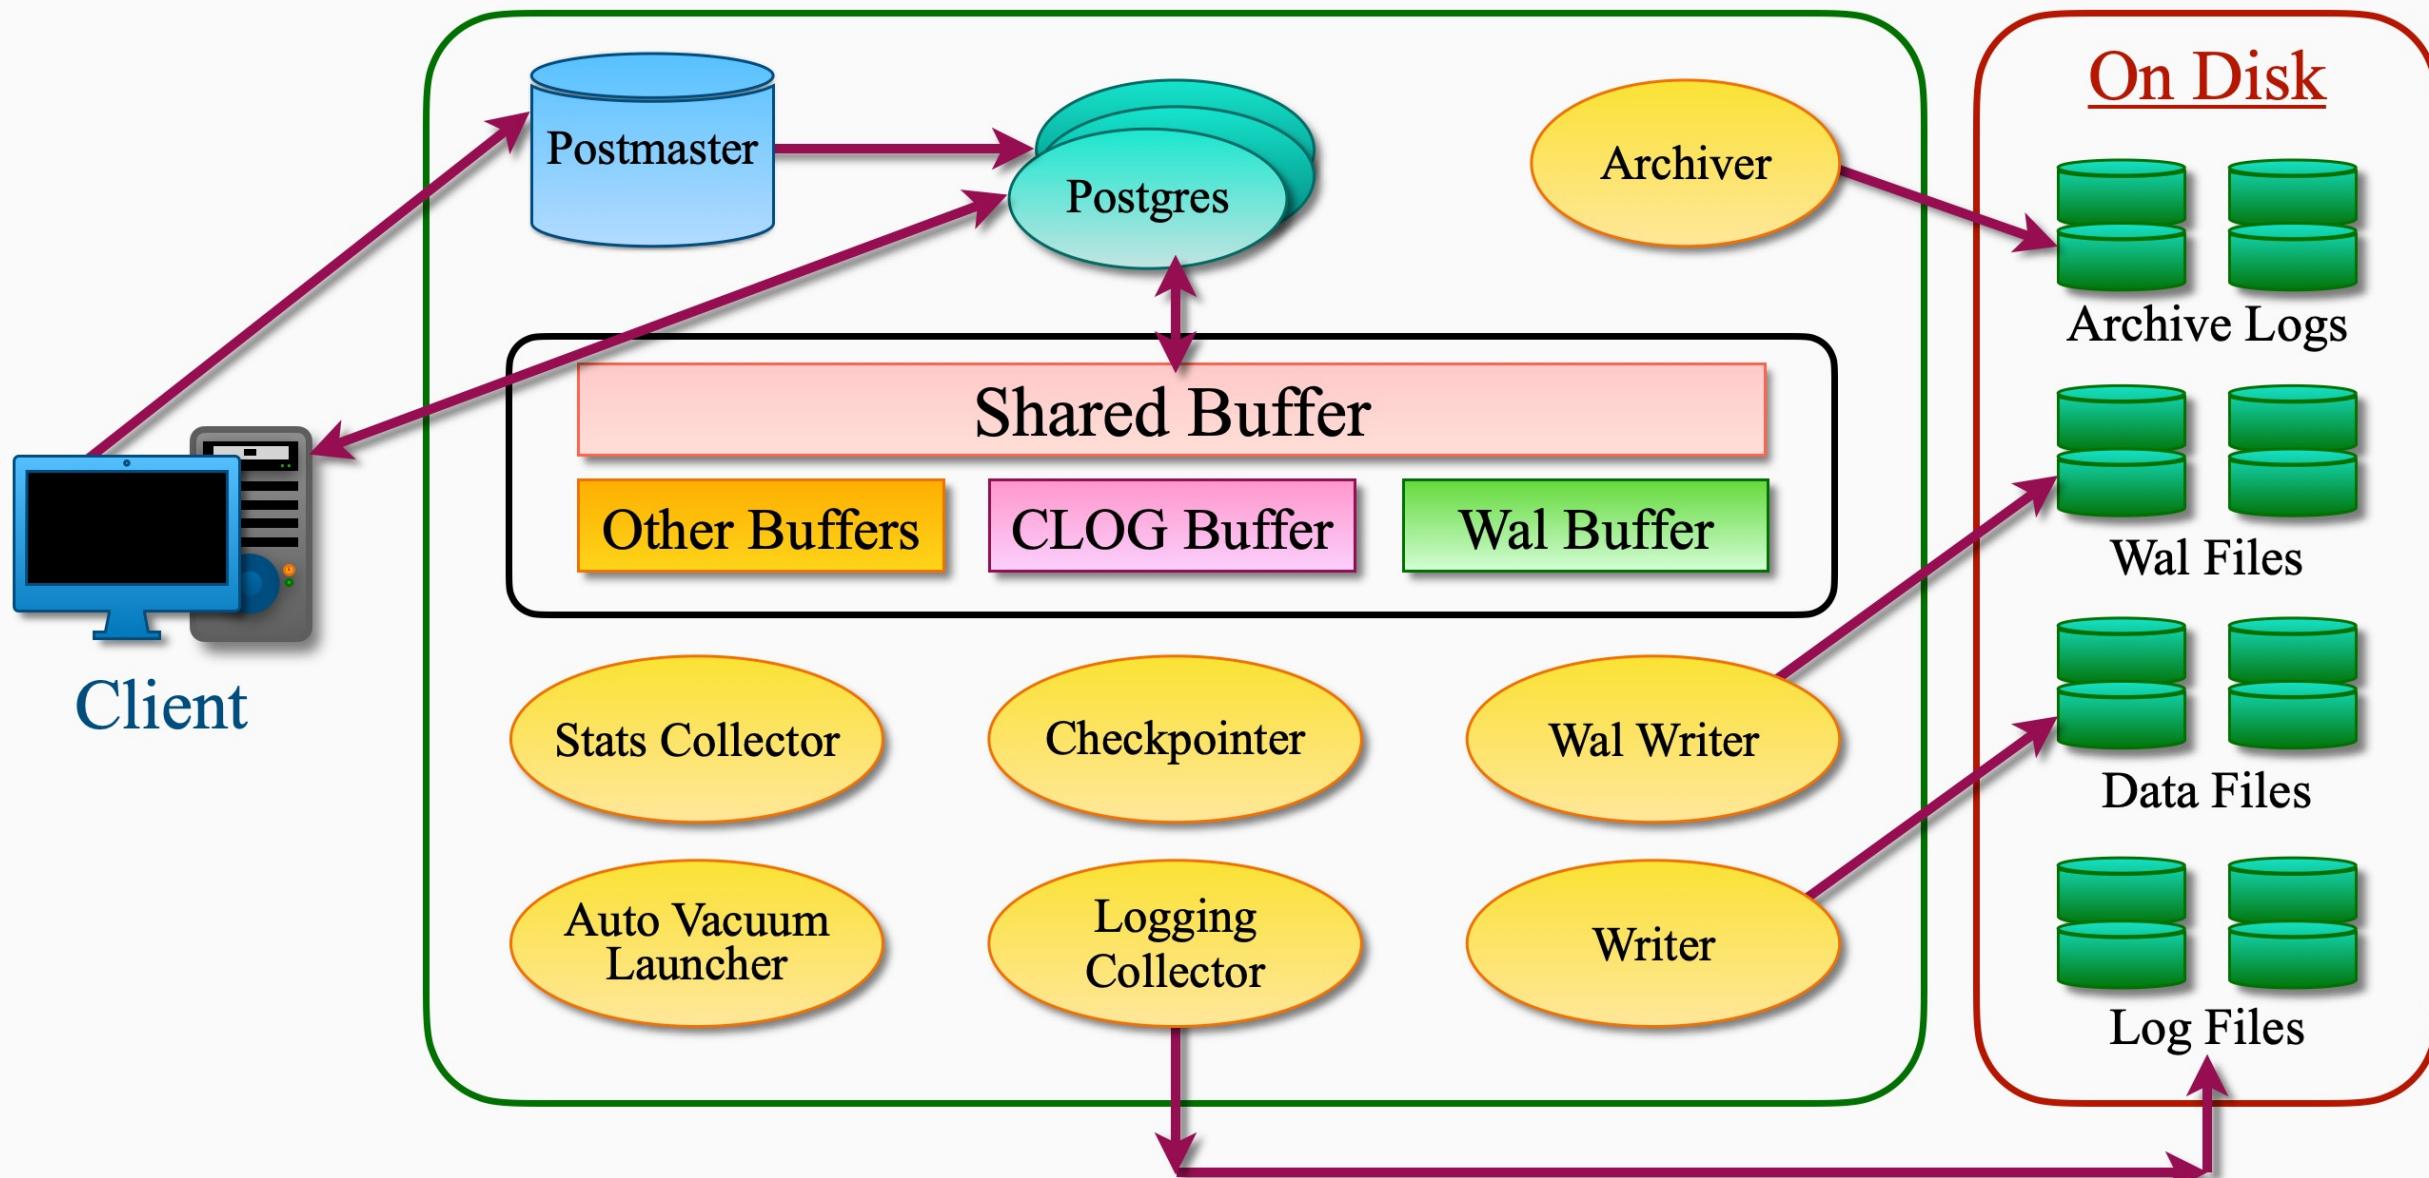


Overview of pg_catalog, schemas, roles, and databases

PostgreSQL vs Oracle: Key architectural differences



PostgreSQL Process and Memory Architecture



PostgreSQL Architecture Key components

1. Postmaster Process (Supervisor)

2. Shared Memory Segments

3. Background Processes

4. Physical File Structure

1. Postmaster Process (Supervisor)

The **first process**

Started when PostgreSQL launches

1. Postmaster Process (Supervisor)

Acts as a **listener**, handles:

Authentication

Authorization

Spawning new backend processes

Monitoring and restarting failed processes

2. Shared Memory Segments

Memory areas used for

query execution,

data caching, and

maintenance tasks.

Shared Buffers



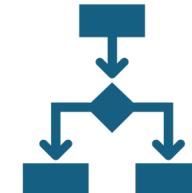
Used for all **read/write operations**



Modified (dirty) pages
are written to disk



by the **background writer**



Controlled by:
`shared_buffers`

WAL Buffers (Write-Ahead Logging)

Temporarily store metadata of changes

Ensures durability and crash recovery

Controlled by: wal_buffers

CLOG Buffers (Commit Logs)



Track **transaction status**



(committed, aborted,
etc.)



Shared across all
sessions

Work Memory

Used for sorting, hashing, joins

Controlled by: `work_mem`

Separate allocation for each operation

Maintenance Work Memory

Used for **index creation,**

foreign key additions, etc.

Controlled by:

`maintenance_work_mem`

Temp Buffers

Used by **temporary tables**

Session-specific

3. Background Processes



Ensure



data consistency



help



maintain database
health.

Background Writer



Writes



dirty pages



from shared
buffers



to disk

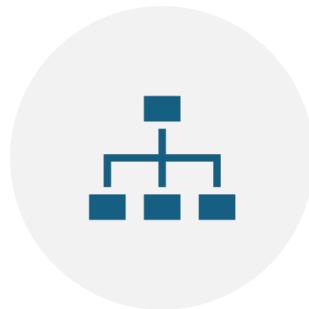
Background Writer



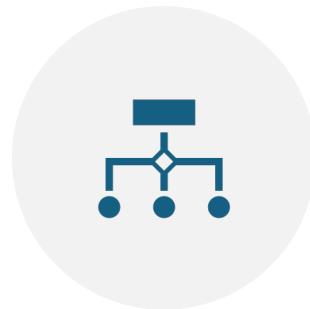
CONTROLS:



BGWRITER_DELAY



BGWRITER_LRU_MAXPAGES



BGWRITER_LRU_MULTIPLIER

Checkpointer

Periodically **flushes**

dirty pages to disk

Essential for

crash recovery

Checkpointer

Triggered by:

Timeout (checkpoint_timeout)

Manual CHECKPOINT command

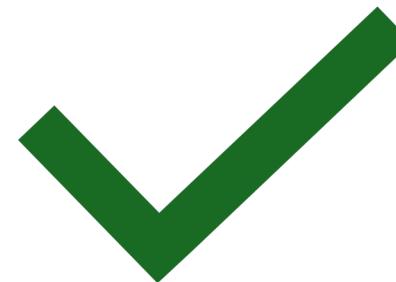
Backup processes

max_wal_size limit

WAL Writer



Writes data from **WAL buffer** to
WAL files on disk



Ensures write operations are
persistent before commit

Autovacuum Launcher



Automates **VACUUM** and
ANALYZE



Maintains **table statistics** and
bloat control

Statistics Collector

Gathers usage metrics:

Query counts

Table scans

Function usage

Statistics Collector

Controlled by:

track_activities

track_functions

track_counts

track_io_timing

Logging Collector

Captures server logs to files

Controlled by:

`logging_collector = ON`

`log_directory = 'pg_log'`

Archiver

Saves WAL segments to archive storage

for point-in-time recovery

Enabled during archive mode

4. Physical File Structure



PostgreSQL uses several file types



to store and manage data persistently:

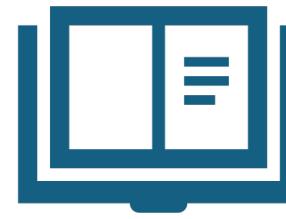
Data Files



Store **actual user data**



Accessed via shared
buffers



for reading/writing

WAL Files



Store **change logs** (Write-Ahead Logging)



Used for crash recovery

Log Files

Store **server logs** including

errors,

warnings, and

execution info

Archive Logs

Archived WAL
segments

for recovery
purposes

Workflow Summary

Client sends a query → received by Postmaster

Postmaster forks a backend (Postgres) process

Data retrieved from Shared Buffers

(or disk if not cached)

Workflow Summary

If write operation:

Change goes to Shared Buffer → WAL Buffer

WAL Writer commits to WAL file

Background Writer/Checkpointer writes to disk

Statistics collected → Logging → Archiving if enabled

Conclusion



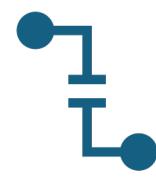
PostgreSQL's
architecture is



robust, scalable,
and



designed for high
reliability,



fault-tolerance,
and concurrency.

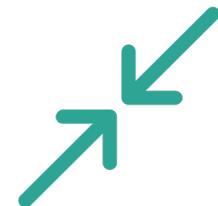
Conclusion



Its modular background
processes and



memory components

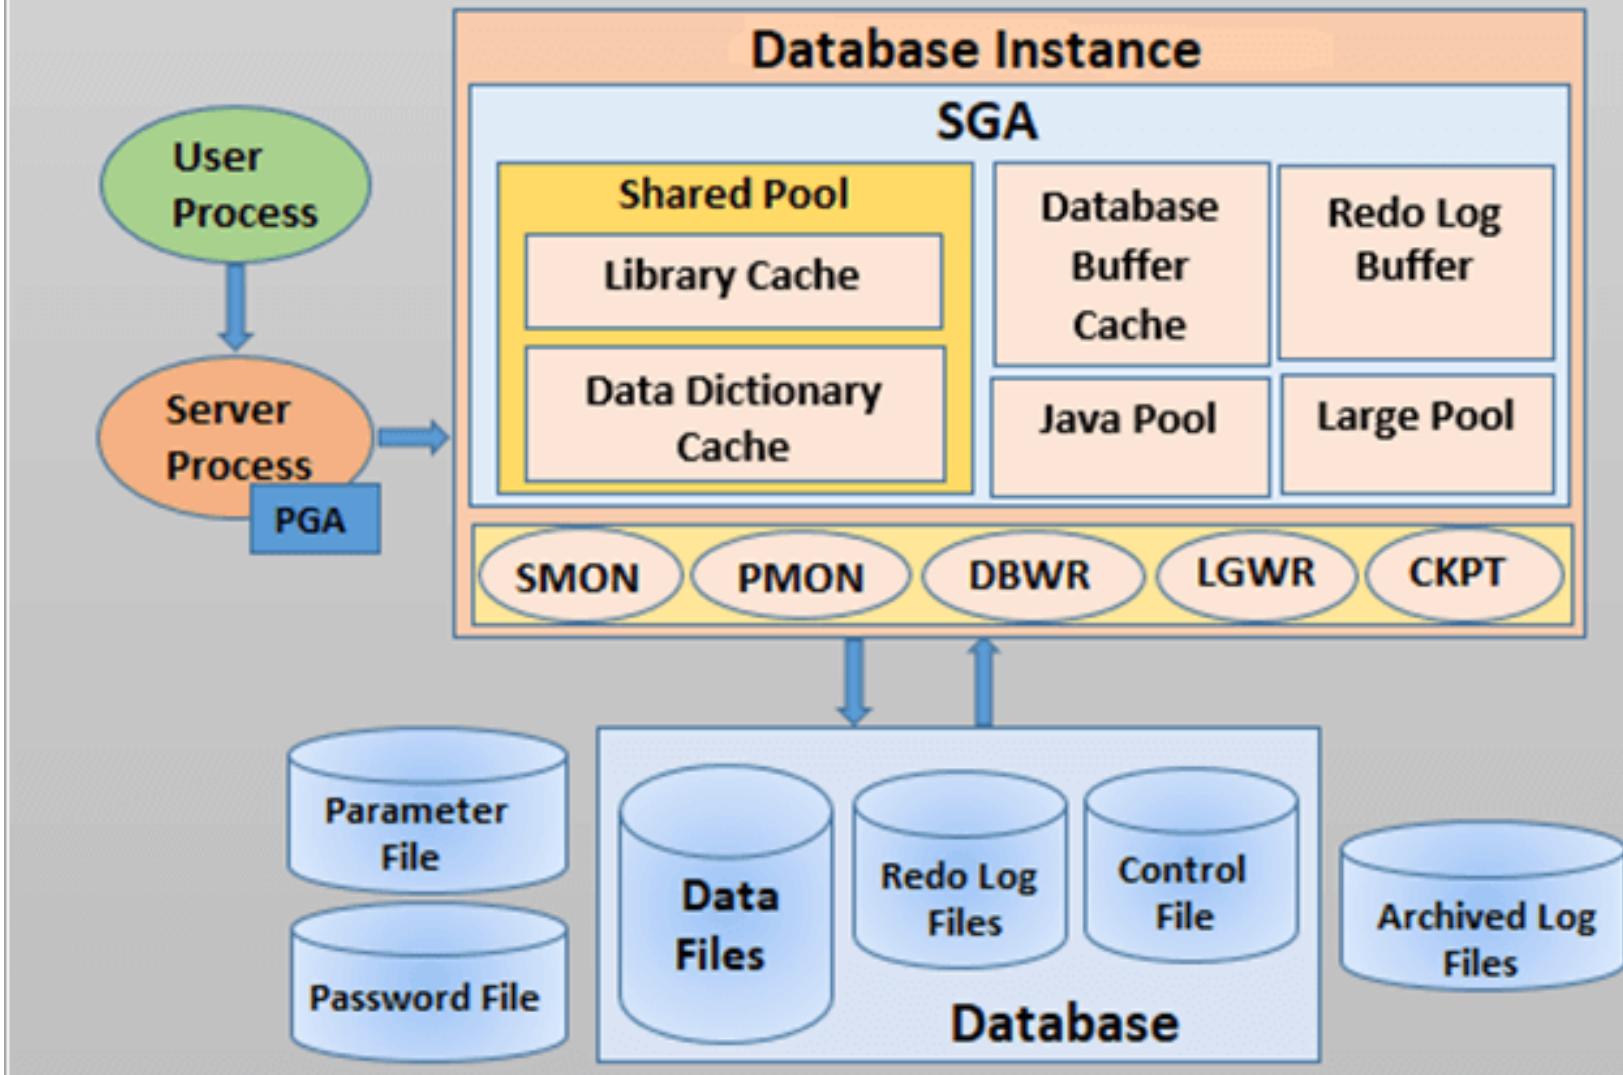


ensure smooth
operations and quick
recovery,



making it a strong choice
for modern applications.

Oracle Database Architecture



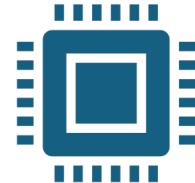
Oracle Database Architecture



Oracle Database is a
multi-model



**relational database
management system**



known for its scalability,
robustness, and



enterprise-grade
features.

Oracle Database Architecture



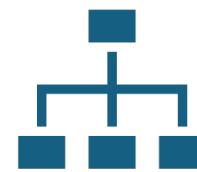
Designed to manage



large amounts of data
efficiently



using a combination of
memory,



processes, files, and
logical storage
structures.

License & Philosophy

Feature	PostgreSQL	Oracle
License	Open Source (PostgreSQL License - similar to BSD)	Commercial (Proprietary License)
Philosophy	Extensibility, Standards Compliance, Community-Driven	Feature-Rich, Enterprise-Grade, Vendor Lock-in

Architecture Overview

Aspect	PostgreSQL	Oracle Database
Process Model	Multi-Process (One process per client connection + background processes)	Multi-Threaded (Single process with multiple threads + background processes)
Memory Architecture	Shared Buffers, Work Mem, WAL Buffers	SGA (System Global Area), PGA (Program Global Area), Redo Log Buffer

Architecture Overview

Aspect	PostgreSQL	Oracle Database
Storage Engine	Single storage engine	Multiple engines (Heap Organized Tables, Index Organized Tables, etc.)
Extensibility	Highly extensible (user-defined types, operators, functions, custom languages)	Limited extensibility (supports PL/SQL, but not new data types easily)

Transaction & Concurrency Control

Feature	PostgreSQL	Oracle
MVCC (Multi-Version Concurrency Control)	Native MVCC (Row-versioning, no read locks for SELECTs)	MVCC via Undo Segments (Undo data in rollback segments)
Locking	Row-level locks, advisory locks	Row-level locks, fine-grained locking, table-level partition locks
Isolation Levels	Read Committed, Repeatable Read, Serializable	Read Committed, Serializable, Read-Only Transactions

Replication & High Availability

Feature	PostgreSQL	Oracle
Replication	Streaming Replication, Logical Replication (built-in), 3rd party (physical), Oracle tools (e.g., Bucardo)	Oracle Data Guard, Oracle GoldenGate (logical)
Sharding	External tools (Citus, pg_shard)	Native (Oracle Sharding)

Data Types & SQL Compliance

Feature	PostgreSQL	Oracle
Data Types	Rich: JSON/JSONB, Arrays, HSTORE, Range Types	Limited: JSON support added in 12c but less flexible
Standards	Very close to ANSI SQL compliance	SQL compliance + proprietary PL/SQL

Partitioning

Feature	PostgreSQL	Oracle
Partitioning	Declarative partitioning (since PostgreSQL 10)	Advanced partitioning (list, range, hash, composite)
Management	Manual in earlier versions; improved automation in recent versions	Fully automated, advanced features

Backup & Recovery

Feature	PostgreSQL	Oracle
Backup	pg_dump (logical), base backups + WAL archiving	RMAN (block-level backup), Data Pump (logical export/import)
Point-in- Time Recovery	WAL-based PITR	Redo/Undo Logs, Flashback Database

PL/SQL vs PL/pgSQL

Feature	PostgreSQL	Oracle
Procedural Language	PL/pgSQL (open), also supports Python, Perl, etc.	PL/SQL (proprietary, tightly integrated)
Triggers/Functions	Easily create functions in multiple languages	Primarily in PL/SQL

Performance & Optimization

Feature	PostgreSQL	Oracle
Optimizer	Cost-based, extensible, adaptive plans in progress	Highly advanced Cost- based Optimizer, adaptive plans, hints
Parallel Query	Supported but less mature	Advanced, with fine- grained control

Security & Auditing

Feature	PostgreSQL	Oracle
Authentication	SSL, GSSAPI, LDAP, Kerberos	Advanced (including fine-grained auditing, TDE, DB Vault)
Auditing	Basic logging and extensions (e.g., pgAudit)	Comprehensive built-in auditing

When to Use What?

PostgreSQL	Oracle
Startups, Open Source Projects	Large Enterprises, Legacy Systems
Cost-sensitive environments	High compliance, mission-critical
Customization & Extensions	Advanced features, support

Summary of Key Differences

PostgreSQL	Oracle
Open-source, community-driven	Proprietary, enterprise-focused
Process-based architecture	Thread-based shared memory architecture
Extensible with custom data types	Fixed feature set, but rich enterprise features
MVCC via row versioning	MVCC via undo segments
Free, with optional paid support	License-based, expensive but feature-rich

Final Thoughts



PostgreSQL: Innovation, Freedom, Extensibility 🚀



Oracle: Stability, Enterprise Support, Advanced Features 🏢



Choose based on your needs, not hype!

PostgreSQL vs Oracle vs SQL Server vs AWS RDS PostgreSQL

Area	PostgreSQL	Oracle	SQL Server	AWS RDS PostgreSQL
Licensing	Free	Very expensive	Expensive	Pay-as-you-use
MVCC	Heap-based	Undo-based	TempDB-based	Same as PostgreSQL
HA	Replication-based	RAC / DG	AlwaysOn	Managed Multi-AZ

PostgreSQL vs Oracle vs SQL Server vs AWS RDS PostgreSQL

Area	PostgreSQL	Oracle	SQL Server	AWS RDS PostgreSQL
OS Access	Full	Full	Full	✗ No
Vacuum / Cleanup	DBA-managed	Undo-managed	Auto	DBA-tuned
Cloud Fit	Excellent	Limited	Moderate	Excellent

Fidelity-Specific Insight

Oracle retained for

core settlement & legacy systems

PostgreSQL used for:

Investment reporting

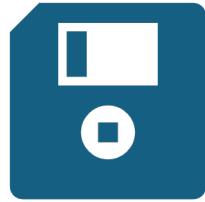
Retirement platforms

Portfolio microservices

Fidelity-Specific Insight



RDS PostgreSQL used
where:



Managed HA &
backups are preferred



OS-level access is not
required

Architecture Differences Local PostgreSQL vs RDS PostgreSQL

Local PostgreSQL (Self-Managed)

DBA controls:

`postgresql.conf`

WAL behavior

Checkpoints

Vacuum tuning

OS, filesystem,
storage

RDS PostgreSQL

AWS
manages:

Instance
lifecycle

Failover

Backups

Patching

RDS PostgreSQL

DBA still
controls:

Schema
design

Indexes

SQL

Vacuum
tuning

Statistics

Query
performance

Key DBA Mindset Shift

RDS removes
sysadmin
work

**Not DBA
responsibility**

Fidelity Use Case



UK / India:



China:



RDS used for regulated production workloads



Managed environments reduce operational risk

Fidelity Use Case

DBAs still troubleshoot:

Slow queries

Blocking

Vacuum lag

Planner misestimates

MVCC – High-Level Concept

Every UPDATE creates a **new row version**

Old versions remain until cleaned by VACUUM

Readers see a consistent snapshot

Writers don't block readers

Oracle vs PostgreSQL

Oracle	PostgreSQL
Undo segments	Heap-based versions
Automatic cleanup	Vacuum required
DBA tunes undo	DBA tunes vacuum

Fidelity Example



PORTFOLIO DASHBOARD
(READ-HEAVY)



TRADE UPDATES (WRITE-
HEAVY)

Fidelity Example

→ MVCC allows:

Real-time dashboards

Without blocking trade execution

⚠ But poor vacuum

silent performance degradation

Important Differences DBAs Must Internalize

No undo tablespaces → **VACUUM is mandatory**

No RAC equivalent → replication-based HA

No optimizer hints (by default)

Planner accuracy depends heavily on statistics

Autovacuum tuning is **non-negotiable**

Typical Oracle DBA Surprise

The database didn't crash

but performance degraded slowly.

That's usually **vacuum + statistics**,

not hardware.

Fidelity-Centric Summary



PostgreSQL is enterprise-ready



DBAs play a **more active performance role**



Same SQL knowledge applies, but **internals differ**



RDS simplifies ops, not tuning



PostgreSQL success = **MVCC + Vacuum + Planner mastery**

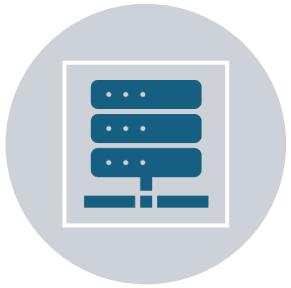
Fidelity Use Case



PostgreSQL used for
investment reporting,



portfolio views,
microservices



Oracle retained for **core**
settlement / legacy
systems



RDS PostgreSQL used
for **managed**
production workloads

Why PostgreSQL Fits These Workloads

PostgreSQL is ideal where workloads are:

Read-heavy

Highly concurrent

Evolving in schema

Cost-sensitive

Cloud-native or microservice-based

Typical Fidelity Workloads

Portfolio dashboards

Investment performance reports

Fund analytics

Customer-facing microservices

Internal analytics platforms

DBA-Relevant Reasons

Requirement	Why PostgreSQL Works
High read concurrency	MVCC ensures readers don't block writers
Frequent updates	No read locks on SELECT
Reporting joins	Strong optimizer + advanced indexing
Schema flexibility	JSONB + partial indexes
Cost control	No license cost
Global access	Works well with replication & read replicas

Tools Overview & Connectivity



psql vs pgAdmin (DBA usage patterns)



Connecting locally



Roles, databases,
schemas



Pagila dataset overview

`psql` vs `pgAdmin`

`psql`

Command Line Interface (CLI)

`pgAdmin 4`

Web-based GUI administration tool

psql (CLI) — DBA Perspective



Lightweight, scriptable,



terminal-based client



Similar to sqlplus
(Oracle) or



sqlcmd (SQL Server)

Why DBAs Prefer psql



Fast during incidents



Scriptable for
automation



Works over SSH / jump
hosts



Minimal overhead

Common DBA Use Cases at Fidelity

Incident troubleshooting (locks, slow queries)

Quick validation during deployments

Running scripts across environments

Connecting to RDS PostgreSQL from bastion hosts

Example

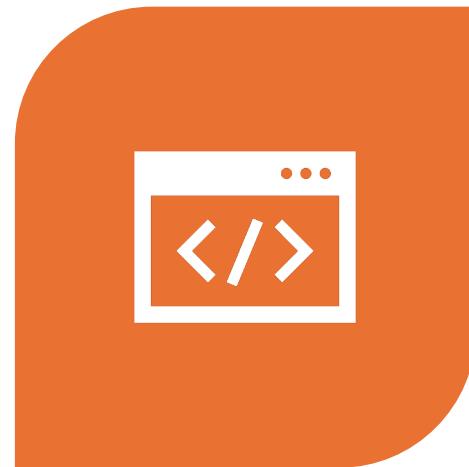
```
psql -h localhost -U postgres
```

```
\l
```

```
\dt
```

```
SELECT * FROM pg_stat_activity;
```

pgAdmin 4 (GUI) — DBA Perspective



WEB-BASED GRAPHICAL
TOOL



SIMILAR TO ORACLE SQL
DEVELOPER / SSMS

Why DBAs Use pgAdmin



Visual query plans



Index and schema
browsing



Explaining query
behavior to teams



Ad-hoc analysis and
learning

Common DBA Use Cases at Fidelity



Query plan analysis for reporting teams



Reviewing index usage



Supporting junior DBAs / developers



Browsing objects in unfamiliar schemas

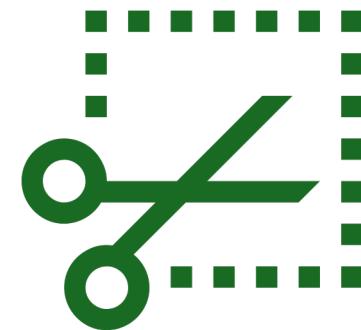
psql vs pgAdmin — Clear Comparison

Area	psql	pgAdmin
Speed		
Automation	Excellent	Limited
Learning Curve	Steep	Easy
Incident Handling	Best choice	Not ideal
Visual Plans	✗	✓
Production Usage	Very common	Selective

DBA Rule of Thumb



psql for operations.



**pgAdmin for analysis and
explanation.**

Connecting Locally to PostgreSQL



Local Connection Model



PostgreSQL supports:



TCP/IP connections



Local socket connections

Connecting Locally to PostgreSQL

- **Typical Local Connection**
- `psql -h localhost -p 5432 -U postgres`

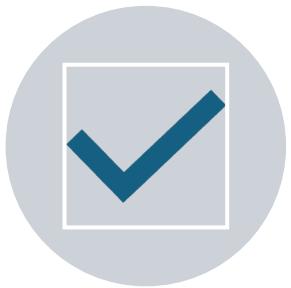
What Happens Internally



Connection request
hits **postmaster**



A backend process is
forked



Session is
authenticated



Backend process
handles all queries

Fidelity Context

Same
connection logic
applies for:

Local
PostgreSQL

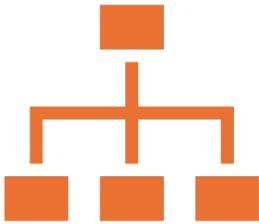
AWS RDS
PostgreSQL

Only difference
in RDS:

SSL enforced

No OS access

Roles, Databases, and Schemas



This hierarchy is
critical for



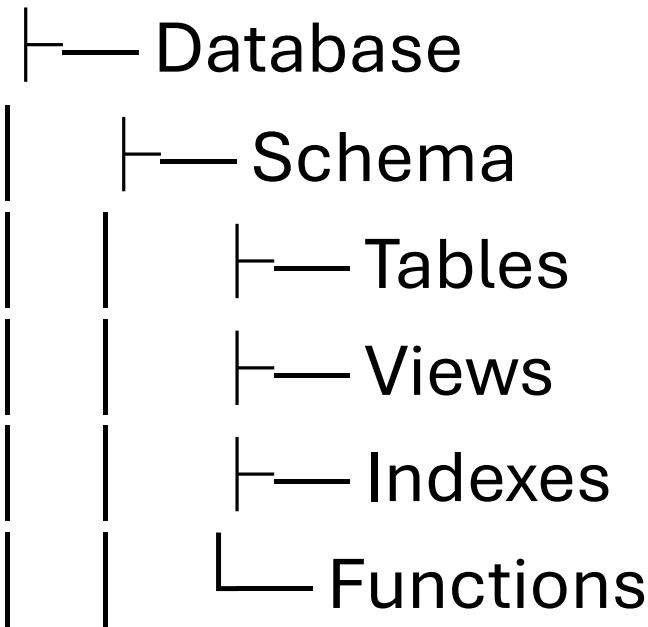
**Oracle / SQL Server
DBAs**



to understand
correctly.

PostgreSQL Object Hierarchy

Cluster



Roles (Users & Groups)

Roles for:

Authentication

Authorization

A role can:

Login (user)

Own objects

Be granted to
other roles
(group)

Oracle Comparison

Oracle	PostgreSQL
User	Role with LOGIN
Role	Role (same concept)

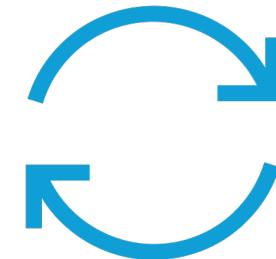
Databases



Logical containers



One connection → one
database



Cannot join across
databases (without FDW)

DBA Impact



Multiple
databases share:



Memory



WAL



Background
processes

Schemas



What Schemas Do



Organize objects



Logical namespaces
inside a database



Control access

Oracle Comparison

Oracle	PostgreSQL
Schema = User	Schema ≠ User

Pagila Dataset Overview

What Pagila Is

Standard PostgreSQL sample database

Equivalent to:

Oracle HR schema

SQL Server AdventureWorks

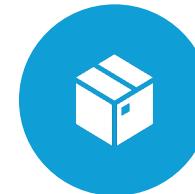
What Pagila Contains



Customers



Transactions



Inventory



Time-based
data



Multiple joins
and indexes

Why Pagila Is Used in This Training

Demonstrates:

Joins

Index usage

Query plans

MVCC
behavior

Vacuum
effects

Fidelity Mapping

Pagila tables
simulate:

Customers →
Investors

Rentals →
Transactions

Payments →
Cash flows

Inventory →
Investment
products

Summary for Fidelity DBA Audience



PostgreSQL tools are **simple but powerful**



psql is the **primary DBA tool**



pgAdmin is a **supporting analysis tool**



Roles ≠ schemas (common Oracle DBA mistake)



Pagila provides realistic, safe practice data

Hands-On Lab

- Connect using psql and pgAdmin
- Create roles, schemas, objects
- Load Pagila schema + data
- Run basic queries

Hands-On Lab

- Measure performance using:
- \timing
- EXPLAIN / EXPLAIN ANALYZE
- Explore system catalogs:
- pg_class
- pg_stat_activity

Connect to PostgreSQL



Database: postgres



Username: postgres



Password: your password



Test Connection & Finish

Summary

Tool	Purpose
PostgreSQL 16	Database Server
pgAdmin 4	Web-based GUI for PostgreSQL
DBeaver	Universal SQL Client & ER Diagram Tool

Useful Links



PostgreSQL Docs: <https://www.postgresql.org/docs/16/>



pgAdmin Docs: <https://www.pgadmin.org/docs/>



DBeaver Docs: <https://dbeaver.io/docs/>

Overview of pg_catalog, schemas, roles, and databases

Surendra Panpaliya

Databases in PostgreSQL



Self-contained unit of data.



Each database has its own



schemas, tables,



users (roles), and configurations

Key Points

PostgreSQL is a multi-database system.

Each database is isolated from others

(no cross-database queries directly).

Default databases



postgres → Default working database



template1 → Template for new databases



template0 → Clean template (minimal)

Schemas



Namespace inside a database.



Think of it as a folder to organize



tables, views, functions, etc.

Key Points



Each database can have **multiple schemas**.



Default schema: public

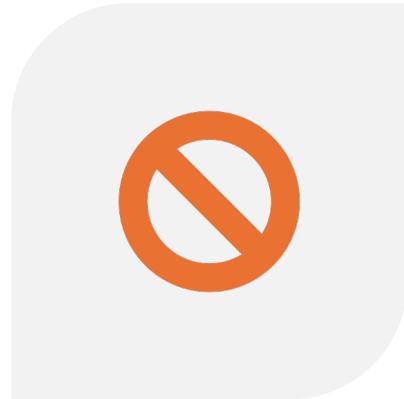


Objects are referenced as:

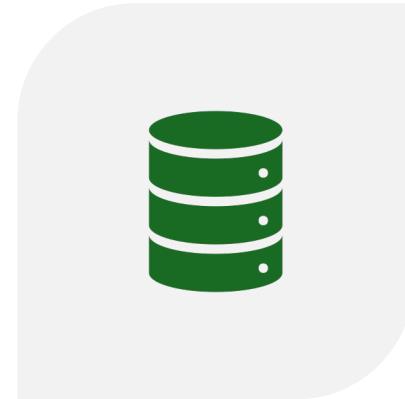


schema_name.object_name

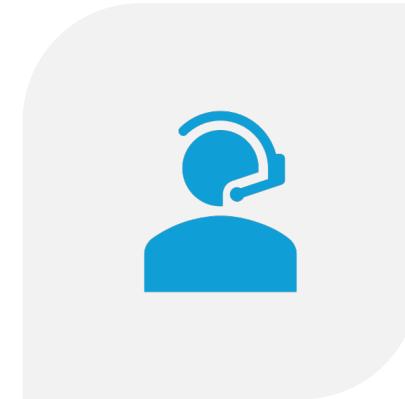
Why Schemas?



AVOID NAME
CONFLICTS



LOGICAL SEPARATION
OF DATA



SUPPORT FOR MULTI-
TENANT APPLICATIONS

Command Examples

-- Create a schema

```
CREATE SCHEMA sales;
```

-- Create table inside schema

```
CREATE TABLE sales.orders (id SERIAL PRIMARY KEY, item TEXT);
```

-- Access table

```
SELECT * FROM sales.orders;
```

pg_catalog



A system schema



Stores PostgreSQL



internal metadata.

Pg_catalog

pg_catalog is a built-in system schema in PostgreSQL

contains all the system tables,

views, data types, functions, and

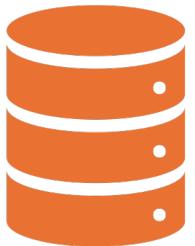
operators used internally by PostgreSQL

to manage the database.

Contains

Object	Purpose
pg_class	Info about tables, indexes, sequences
pg_user	View of roles with login
pg_roles	Info about all roles
pg_namespace	Info about schemas
pg_tables	List of tables
pg_indexes	List of indexes

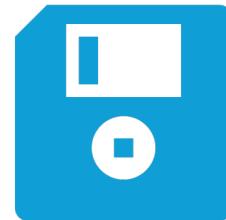
Use Case:



Query database
metadata



Understand internal
PostgreSQL structures



System operations,
backups, monitoring

Example Queries:

-- List all tables using pg_catalog

```
SELECT tablename FROM pg_catalog.pg_tables WHERE  
schemaname = 'public';
```

-- List all roles

```
SELECT rolname FROM pg_catalog.pg_roles;
```

Roles

A Role in PostgreSQL = User Account
+ Group

Roles manage authentication &
authorization.

Types of Roles:

Role	Description
Login Role	Can log in (CREATE ROLE username LOGIN;)
Group Role	Used for permissions, no login

Privileges Controlled By Roles:

CONNECT	CONNECT (to database)
CREATE	CREATE (schemas, tables)
SELECT	SELECT, INSERT, UPDATE, DELETE (on tables)
EXECUTE	EXECUTE (functions)

Role Management Examples



-- Create role with login



```
CREATE ROLE devuser LOGIN PASSWORD 'mypassword';
```



-- Grant role permissions



```
GRANT CREATE ON DATABASE mydb TO devuser;
```

Role Management Examples

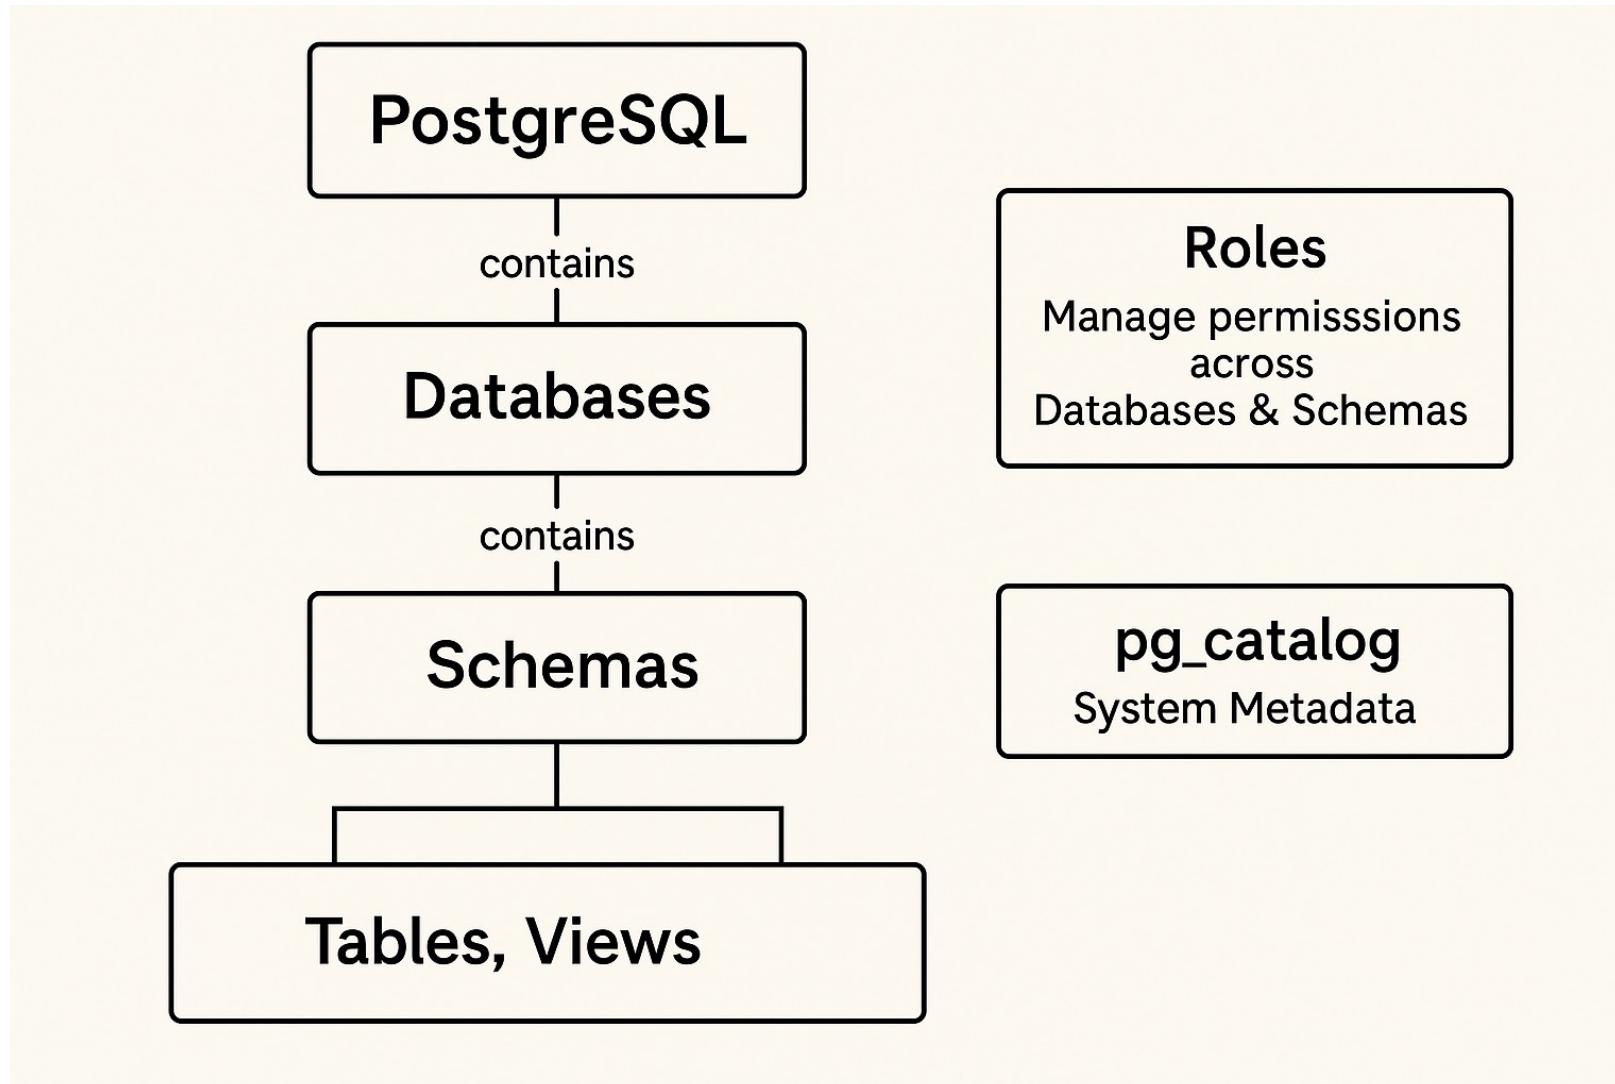
-- Create group role

```
CREATE ROLE analyst;
```

-- Assign user to group

```
GRANT analyst TO devuser;
```

Relationship Diagram



Summary Table

Concept	Purpose
Database	Top-level data container
Schema	Namespace within a database
<code>pg_catalog</code>	Stores metadata about objects
Role	User and permission management

Real-World Analogy

PostgreSQL Concept	Analogy
Database	Apartment Building
Schema	Individual Apartment
Table/View	Furniture inside an apartment
Role	Security Guard controlling access
pg_catalog	Building Management Records



**Thank you for
your support and
patience**

Surendra Panpaliya
Founder and CEO
GKTCS Innovations
<https://www.gktcs.com>